Explaining how a prediction was made is increasingly becoming a necessity rather than just an option. But machine learning models are complex, often leaving users in the dark about how they make predictions.

Enter SHAP - an essential addition to any data scientist's toolkit. With SHAP, you can analyze how each feature contributes to a prediction. This Explainable AI technique helps you understand not just individual predictions, but the model's overall behavior.

SHAP works for any machine learning model - whether it's a linear model or a deep neural network - and for many types of data, from tabular data to images.

This book:

- Breaks down the intuition behind SHAP.
- Gives you the practical know-how to apply SHAP to real-world tasks.
- Dives into the theory behind Shapley values and SHAP.
- Provides hands-on Python code examples.

Whether you're a seasoned data scientist or a newcomer to the field, this book is an invaluable resource that will undoubtedly deepen your understanding of SHAP values and their significance in the realm of interpretable machine learning.

– Junaid Butt, Research Software Engineer, IBM Research

This book takes readers on a comprehensive journey from foundational concepts to practical applications of SHAP

– Carlos Mougan, Marie Skłodowska-Curie AI Ethics Researcher.

INTERPRETING MACHINE LEARNING MODELS WITH SHAP

CHRISTOPH MOLNAR

Interpreting Machine Learning Models With SHAP

A Guide With Python Examples And Theory On Shapley Values

Christoph Molnar